

Recipe Similarity Using Ingredient Extraction and Collaborative Filtering

Benjamin Holtz
Department of Computer Science
Stanford University
Stanford, CA 94305
bholtz@stanford.edu

Xiangnong Wang
Public Policy Program
Stanford University
Stanford, CA 94305
gxwang@stanford.edu

Abstract

In this report, we present an approach to determining similarity between recipes using information extraction and collaborative filtering techniques. Recipe similarity can be used to gain a functional understanding of what distinguishes different types of cuisine. Leveraging recipe data from the WikiBooks Cookbook¹, we can identify natural recipe classes using high-level user generated tags (e.g. "Indian Cuisine"). This report will describe various approaches to improving upon a baseline bag-of-words model.

1 Notes

This report was developed as part of the final project for Stanford University's CS 224U: Natural Language Understanding taught by Bill MacCartney and Chris Potts. The source code can be downloaded from the project's GitHub page².

2 Introduction

Food often characterizes culture, and indeed, some cultures characterize themselves by their food. Anecdotally, we can make conjectures about cuisines that share close ties, but to what extent do recipes actually differ between these classifications? Thus, this report describes a system to compare recipe similarities within classes and across classes. We use the WikiBooks Cookbook data because it is freely available and is meta-tagged by users based

on recipe characteristics. Some examples of recipe-wide tags are food categories like "bread recipes" and "fried recipes," as well as, cuisines like "Indian recipes" and "Chinese recipes." Thus, the WikiBooks Cookbook dataset includes both documents and the natural classes that include these documents. Recipes are interesting from a similarity metric point-of-view, because there may be some functional reasons for similarity and dis-similarity. For instance, bread recipes likely need flour and yeast. On the other hand, while Indian recipes and Mexican recipes both use spices, the cuisines have different flavor profiles precisely due to the specific spices that are commonly used.

The functional contribution of different ingredients to the flavor profile of a recipe highlights the significance of information extraction for the recipe similarity task. While bag-of-words alone can provide a useful baseline similarity score between recipes and class, information extraction on ingredients lists can capture more fine-grained information about the ingredient itself. Ingredients lists are generally short semi-structured phrases or fragments that include a base ingredient (e.g. onion), a unit of measurement and its quantity (e.g. 2 cups), and some properties (e.g. diced). However, simply running these lines through a parser doesn't provide a clean and consistent way of extracting these data, so there is plenty of room to improve extraction through hand-built features. Information extraction on ingredients lists can tease out functional information about the ingredients and their quantities.

¹<http://en.wikibooks.org/wiki/Cookbook>

²http://github.com/gxwang/cs224u_recipe

3 Methodology

3.1 Data Processing

Our data comes from Wikimedia as a single file of 2,403 XML-formatted recipes. The file is only formatted in a way that reflects its status as a set of Wikimedia pages and not as a set of recipes. Thus, there is substantial pre-processing involved in re-formatting recipes from XML to recipe objects. Immediately, we can extract any XML flagged information like the title of the recipe and the useful recipe category tags associated with each page. However, the actual text of a Cookbook recipe is all included in a single XML flag, and thus does not have any XML formatting. Fortunately, there is some structure that is gained from typical Wiki page internal formatting. For example, ingredients lists are often notated with text headings like "==Ingredients==." We can rely on the fact that Wikimedia contributors keep WikiBooks Cookbook pages reasonably consistent and cleanly formatted. Thus, we can reliably strip away the unnecessary data and keep only the data we need. The output of the pre-processing is a set of recipe objects with ingredient lines that can be parsed to extract the ingredients from these lines.

3.2 Ingredient Extraction

There are three main tasks for ingredient extraction: identifying units and quantities, identifying base ingredients, and identifying ingredient properties. For each ingredient line, our goal is to extract each of those pieces of data. Ultimately, we achieve this through a set of hand-built features to extract this information, which works to a varying degree of success.

First, to identify the units of measurement, we put together a relatively exhaustive list of cooking units (e.g. cups, liters, cloves) and all common variations, expressing them as a long regular expression statement. Running an ingredient line through the Stanford English PCFG parser and examining the collapsed dependencies, we note unit and quantity pairs by identifying a unit using the regular expression described above and the "num" or "number" relations. We further attempt to normalize the units by converting fractional quantities into their decimal form and all liquid and weight units to a base measurement.

For example, 1/2 cups is converted to 118.2941 mL, and 1 lb is converted to 453.5924 grams.

Next, to determine the base ingredient in an ingredient line, we perform a second run of the Stanford parser. This time, we remove the already extracted unit and quantity, as well as the word "of," because they often interfere with identifying the base ingredient. Then, using the "root" relation from the collapsed dependency parse, we can extract base ingredient. If the "root" relation fails to find a noun, we look for the base ingredient when iterating through the other relations.

Finally, to identify the properties of an ingredient, we use the "nn," "appos," "amod," and "nsubj" relations from the second dependency parse. For example, the line "1 cup chopped onions" returns "onions" as the base ingredient and "chopped" as the property. Our basic approach ignores some things like compound ingredients. So, 3 teaspoons baking powder would lead to powder as the base ingredient and "baking" as its property.

3.3 Electronically-available resources

ACL2012 provides this description in L^AT_EX2e (acl2012.tex) and PDF format (acl2012.pdf), along with the L^AT_EX2e style file used to format it (acl2012.sty) and an ACL bibliography style (acl.bst). These files are all available at <http://www.acl2012.org>. A Microsoft Word template file (acl2012.dot) is also available at the same URL. We strongly recommend the use of these style files, which have been appropriately tailored for the ACL2012 proceedings. If you have an option, we recommend that you use the L^AT_EX2e version. **If you will be using the Microsoft Word template, we suggest that you anonymize your source file so that the pdf produced does not retain your identity.** This can be done by removing any personal information from your source document properties.

3.4 Format of Electronic Manuscript

For the production of the electronic manuscript you must use Adobe's Portable Document Format (PDF). This format can be generated from postscript files: on Linux/Unix systems, you can use `ps2pdf` for this purpose; under Microsoft Windows, you can use Adobe's Distiller, or if you have `cygwin` installed, you can use `dvipdf` or `ps2pdf`. Note

that some word processing programs generate PDF which may not include all the necessary fonts (esp. tree diagrams, symbols). When you print or create the PDF file, there is usually an option in your printer setup to include none, all or just non-standard fonts. Please make sure that you select the option of including ALL the fonts. *Before sending it, test your PDF by printing it from a computer different from the one where it was created.* Moreover, some word processor may generate very large postscript/PDF files, where each page is rendered as an image. Such images may reproduce poorly. In this case, try alternative ways to obtain the postscript and/or PDF. One way on some systems is to install a driver for a postscript printer, send your document to the printer specifying “Output to a file”, then convert the file to PDF.

Additionally, it is of utmost importance to specify the **US-Letter format** (8.5in \times 11in) when formatting the paper. When working with dvips, for instance, one should specify `-t letter`.

Print-outs of the PDF file on US-Letter paper should be identical to the hardcopy version. If you cannot meet the above requirements about the production of your electronic submission, please contact the publication chair above as soon as possible.

3.5 Layout

Format manuscripts two columns to a page, in the manner these instructions are formatted. The exact dimensions for a page on US-letter paper are:

- Left and right margins: 1in
- Top margin: 1in
- Bottom margin: 1in
- Column width: 3.15in
- Column height: 9in
- Gap between columns: 0.2in

Papers should not be submitted on any other paper size. If you cannot meet the above requirements about the production of your electronic submission, please contact the publication chair above as soon as possible.

3.6 Fonts

For reasons of uniformity, Adobe’s **Times Roman** font should be used. In L^AT_EX2e this is accomplished by putting

```
\usepackage{times}
\usepackage{latexsym}
```

in the preamble. If Times Roman is unavailable, use **Computer Modern Roman** (L^AT_EX2e’s default). Note that the latter is about 10% less dense than Adobe’s Times Roman font.

Type of Text	Font Size	Style
paper title	15 pt	bold
author names	12 pt	bold
author affiliation	12 pt	
the word “Abstract”	12 pt	bold
section titles	12 pt	bold
document text	11 pt	
captions	11 pt	
abstract text	10 pt	
bibliography	10 pt	
footnotes	9 pt	

Table 1: Font guide.

3.7 The First Page

Center the title, author’s name(s) and affiliation(s) across both columns. Do not use footnotes for affiliations. Do not include the paper ID number assigned during the submission process. Use the two-column format only when you begin the abstract.

Title: Place the title centered at the top of the first page, in a 15 point bold font. (For a complete guide to font sizes and styles, see Table 1.) Long title should be typed on two lines without a blank line intervening. Approximately, put the title at 1in from the top of the page, followed by a blank line, then the author’s names(s), and the affiliation on the following line. Do not use only initials for given names (middle initials are allowed). Do not format surnames in all capitals (e.g., “Zhou,” not “ZHOU”). The affiliation should contain the author’s complete address, and if possible an electronic mail address. Leave about 0.75in between the affiliation and the body of the first page. The title, author names and addresses should be completely identical to those

entered to the electronic paper submission website in order to maintain the consistency of author information among all publications of the conference.

Abstract: Type the abstract at the beginning of the first column. The width of the abstract text should be smaller than the width of the columns for the text in the body of the paper by about 0.25in on each side. Center the word **Abstract** in a 12 point bold font above the body of the abstract. The abstract should be a concise summary of the general thesis and conclusions of the paper. It should be no longer than 200 words. The abstract text should be in 10 point font.

Text: Begin typing the main body of the text immediately after the abstract, observing the two-column format as shown in the present document. Do not include page numbers.

Indent when starting a new paragraph. For reasons of uniformity, use Adobe's **Times Roman** fonts, with 11 points for text and subsection headings, 12 points for section headings and 15 points for the title. If Times Roman is unavailable, use **Computer Modern Roman** (L^AT_EX2e's default; see section 3.4 above). Note that the latter is about 10% less dense than Adobe's Times Roman font.

3.8 Sections

Headings: Type and label section and subsection headings in the style shown on the present document. Use numbered sections (Arabic numerals) in order to facilitate cross references. Number subsections with the section number and the subsection number separated by a dot, in Arabic numerals. Do not number subsubsections.

Citations: Citations within the text appear in parentheses as (Gusfield, 1997) or, if the author's name appears in the text itself, as Gusfield (1997). Append lowercase letters to the year in cases of ambiguities. Treat double authors as in (Aho and Ullman, 1972), but write as in (Chandra et al., 1981) when more than two authors are involved. Collapse multiple citations as in (Gusfield, 1997; Aho and Ullman, 1972). Also refrain from using full citations as sentence constituents. We suggest that instead of

“(Gusfield, 1997) showed that ...”

you use

“Gusfield (1997) showed that ...”

If you are using the provided L^AT_EX and BibT_EX style files, you can use the command `\newcite` to get “author (year)” citations.

As reviewing will be double-blind, the submitted version of the papers should not include the authors' names and affiliations. Furthermore, self-references that reveal the author's identity, e.g.,

“We previously showed (Gusfield, 1997)
...”

should be avoided. Instead, use citations such as

“Gusfield (1997) previously showed ... ”

Please do not use anonymous citations and do not include acknowledgements when submitting your papers. Papers that do not conform to these requirements may be rejected without review.

References: Gather the full set of references together under the heading **References**; place the section before any Appendices, unless they contain references. Arrange the references alphabetically by first author, rather than by order of occurrence in the text. Provide as complete a citation as possible, using a consistent format, such as the one for *Computational Linguistics* or the one in the *Publication Manual of the American Psychological Association* (American Psychological Association, 1983). Use of full names for authors rather than initials is preferred. A list of abbreviations for common computer science journals can be found in the *ACM Computing Reviews* (Association for Computing Machinery, 1983).

The L^AT_EX and BibT_EX style files provided roughly fit the American Psychological Association format, allowing regular citations, short citations and multiple citations as described above.

Appendices: Appendices, if any, directly follow the text and the references (but see above). Letter them in sequence and provide an informative title: **Appendix A. Title of Appendix.**

Acknowledgment sections should go as a last section immediately before the references. Do not number the acknowledgement section.

3.9 Footnotes

Footnotes: Put footnotes at the bottom of the page. They may be numbered or referred to by asterisks or other symbols.³ Footnotes should be separated from the text by a line.⁴ Footnotes should be in 9 point font.

3.10 Graphics

Illustrations: Place figures, tables, and photographs in the paper near where they are first discussed, rather than at the end, if possible. Wide illustrations may run across both columns and should be placed at the top of a page. Color illustrations are discouraged, unless you have verified that they will be understandable when printed in black ink.

Captions: Provide a caption for every illustration; number each one sequentially in the form: “Figure 1. Caption of the Figure.” “Table 1. Caption of the Table.” Type the captions of the figures and tables below the body, using 10 point text.

4 Translation of non-English Terms

It is also advised to supplement non-English characters and terms with appropriate transliterations and/or translations since not all readers understand all such characters and terms.

Inline transliteration or translation can be represented in the order of: original-form transliteration “translation”.

5 Length of Submission

Long papers may consist of up to eight (8) pages of content and an additional two (2) pages of references, and short papers may consist of up to four (4) pages of content and two (2) additional pages of references. Papers that do not conform to the specified length and formatting requirements are subject to re-submission.

6 Other Issues

Those papers that had software and/or dataset submitted for the review process, should also submit it with the camera-ready paper. Besides, the software and/or dataset should not be anonymous.

³This is how a footnote should appear.

⁴Note the line separating the footnotes from the text.

Please note that the publications of ACL2012 will be publicly available at ACL Anthology (<http://aclweb.org/anthology-new/>) on July 2nd, 2012, one week before the start of the conference. Since some of the authors may have plans to file patents related to their papers in the conference, we are sending this reminder that July 2nd, 2012 may be considered to be the official publication date, instead of the opening day of the conference.

Acknowledgments

Do not number the acknowledgment section. Do not include this section when submitting your paper for review.

References

- Alfred V. Aho and Jeffrey D. Ullman. 1972. *The Theory of Parsing, Translation and Compiling*, volume 1. Prentice-Hall, Englewood Cliffs, NJ.
- American Psychological Association. 1983. *Publications Manual*. American Psychological Association, Washington, DC.
- Association for Computing Machinery. 1983. *Computing Reviews*, 24(11):503–512.
- Ashok K. Chandra, Dexter C. Kozen, and Larry J. Stockmeyer. 1981. Alternation. *Journal of the Association for Computing Machinery*, 28(1):114–133.
- Dan Gusfield. 1997. *Algorithms on Strings, Trees and Sequences*. Cambridge University Press, Cambridge, UK.